

Peer review of group: tortellini-ravioli

By group: running_late1

Code and materials review

The project has a clear and useful goal: building a chatbot that answers questions based on Grafana documentation. This is a suitable use case for a RAG-based system, since Grafana documentation is large and often difficult to search manually, as stated in the report.

The use of Ollama for local inference is also a good choice, because it avoids dependence on paid external APIs.

However, the project is not very easy to reproduce. Running it required several manual steps that were either missing or not clearly explained, so after an hour or so of trying to run it I couldn't get the project to start (the other reviewer had the same problem as well). The system expects Grafana documentation to be available in a certain file, but this data was not already present after downloading the project. Some Python dependencies also appeared to be missing, such as pyyaml and pandas. This suggests that the setup instructions or requirements.txt are incomplete.

The program also does not give very user-friendly messages when required folders or data are missing.

Overall, the implementation is conceptually good, but the setup process needs improvement. A reviewer should be able to run the project with fewer manual corrections if any at all.

Report review - Good points

The project addresses a realistic and everyday problem and uses an appropriate solution for it. The report explains the motivation and general architecture clearly. It is understandable why a chatbot for Grafana documentation is useful, and the report gives a reasonable overview.

Using local models through Ollama is a positive point, especially from the perspective of privacy and avoiding external API costs.

The evaluation is easy to understand. It contains a relatively small number of questions, but the scoring appears to be mostly manual and subjective. It would be stronger if the authors used clearer objective evaluation criteria.

The project also includes a benchmark script, which shows that the authors attempted to evaluate the system instead of only presenting the implementation. Including problematic cases shows that the authors paid attention to detail and the model analysis in Table 2 is a really nice touch to finish off with.

Report review - Bad points

The main weakness is reproducibility. I could not get the project to run despite following all the instructions. It required manual downloading of Grafana documentation, manual installation of missing dependencies, and additional troubleshooting.

The documentation should clearly state all required steps: installing Ollama, pulling the required models, preparing the Grafana documentation folder, installing dependencies, building the database, and only then running the query scrip, especially because the practical installation process is the weakest part of the project.

Overall score: 8 (good)