

Peer review of group: neznaniletecipredmet

By group: running_late1

The project aims to build a retrieval-augmented chatbot for UL FRI students that can answer questions based on official university and faculty web content. This is a meaningful application of RAG systems, since university information is can usually be found across many webpages and is often difficult to search manually.

The overall system design is solid and follows a modern RAG pipeline: crawling, parsing, chunking, embedding, FAISS indexing, and retrieval with optional reranking.

The project structure is also well organized, with a clear separation between scripts, source code, and evaluation artefacts. The inclusion of a prebuilt FAISS index and evaluation results makes it easier to inspect the system without re-running the full pipeline, which is very useful for grading purposes.

The evaluation methodology is also a strong point. The project does not rely only on qualitative examples, but includes systematic evaluation with a curated question set and multiple model comparisons. The inclusion of LLM-as-a-judge evaluation is especially interesting and aligns with current research practices in RAG evaluation.

The report also clearly documents different experimental configurations (embeddings, chunk sizes, retrieval modes), which helps to understand how final conclusions were reached.

In terms of output quality, the system generally produces relevant and context-aware answers when the correct documents are retrieved. The responses are usually grounded in the provided university data and avoid obvious hallucinations. However, the quality is strongly dependent on retrieval performance: in cases where relevant text is not included in the retrieved documents, the model may produce incomplete or overly generic answers.

Furthermoer, reproducibility is somewhat challenging in practice. While the repository contains a large amount of code and data, running the full system required careful environment setup and dependency handling. In particular, the retrieval pipeline failed initially due to environment-related issues (embedding model download and disk space issues, since you had to have 1.1GB of space).

Additionally, although the README is quite detailed, a first-time user may still struggle with deciding the exact minimal steps needed to run only a basic query successfully without triggering the full pipeline or additional evaluation components.

Finally, it would be nice if there were a “quick run” mode that guarantees a working retrieval output with minimal setup. This would make the project more accessible.

Overall score: 8 (good)

